

INDEX

Sr.no.	Practical	Date	Sign
1.	<p>Data Pre-processing and Exploration</p> <ol style="list-style-type: none"> Load a CSV dataset. Handle missing values, inconsistent formatting, and outliers. Load a dataset, calculate descriptive summary statistics, create visualizations using different graphs, and identify potential features and target variables Note: Explore Univariate and Bivariate graphs (Matplotlib) and Seaborn for visualization Create or Explore datasets to use all pre-processing routines like label encoding, scaling, and binarization. 		
2.	<p>Testing Hypothesis</p> <ol style="list-style-type: none"> Implement and demonstrate the FIND-S algorithm for finding the most specific hypothesis based on a given set of training data samples. Read the training data from CSV file and generate the final specific hypothesis. (Create your dataset) 		
3.	<p>Linear Models</p> <ol style="list-style-type: none"> Simple Linear Regression Fit a linear regression model on a dataset. Interpret coefficients, make predictions, and evaluate performance using metrics like R-squared and MSE Multiple Linear Regression Extend linear regression to multiple feature. Handle feature selection and potential multicollinearity Regularized Linear Models Implement Regression variants like LASSO and Ridge on any generated dataset 		
4.	<p>Discriminative Models</p> <ol style="list-style-type: none"> Logistic Regression : Perform binary classification using logistic regression. Calculate accuracy, precision, recall, and understand the ROC curve." Implement and demonstrate k-nearest Neighbor algorithm. Read the training data from a .CSV file and build the model to classify a test sample. Print both correct and wrong predictions. Build a decision tree classifier or regressor. Control hyperparameters like tree depth to avoid overfitting. Visualize the tree. Implement a Support Vector Machine for any relevant dataset. Train a random forest ensemble. Experiment with the number of trees and feature sampling. Compare performance to a single decision tree. Implement a gradient boosting machine (e.g., XGBoost). Tune hyperparameters and explore feature importance. 		

5.	Generative Models <ul style="list-style-type: none"> a. Implement and demonstrate the working of a Naive Bayesian classifier using a sample data set. Build the model to classify a test sample. b. Implement Hidden Markov Models using hmmlearn 		
6.	Probabilistic Models <ul style="list-style-type: none"> a. Implement Bayesian Linear Regression to explore prior and posterior distribution. b. Implement Gaussian Mixture Models for density estimation and unsupervised clustering. 		
7.	Model Evaluation and Hyperparameter Tuning <ul style="list-style-type: none"> a. Implement cross-validation techniques (k-fold, stratified, etc.) for robust model evaluation b. Systematically explore combinations of hyperparameters to optimize model performance.(use grid and randomized search) 		
8.	Bayesian Learning <ul style="list-style-type: none"> a. Implement Bayesian Learning using inferences 		

Practical 1: Data Pre-processing and Exploration

1a. Load a CSV dataset. Handle missing values, inconsistent formatting, and outliers.

1. Import Libraries

Import necessary libraries

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

2. Load the Dataset

Load the Titanic dataset from a URL

```
url = "https://raw.githubusercontent.com/datasciencedojo/datasets/master/titanic.csv" data
= pd.read_csv(url)
```

Display the first few rows

```
print(data.head())
```

3. Handle Missing Values

Check for missing values

```
print("Missing values in each column:")
print(data.isnull().sum())
```

Fill missing values in 'Age' with the mean

```
data['Age'].fillna(data['Age'].mean(), inplace=True)
```

Fill missing values in 'Embarked' with the most common value

```
data['Embarked'].fillna(data['Embarked'].mode()[0], inplace=True)
```

Drop rows where 'Cabin' is missing (too many NaNs)

```
data.drop(columns=['Cabin'], inplace=True)
```

Verify missing values are handled

```
print("\nAfter handling missing values:")
print(data.isnull().sum())
```

4. Fix Inconsistent Formatting

Fix inconsistent formatting in the 'Sex' column

```
data['Sex'] = data['Sex'].str.lower().str.strip()
```

Verify unique values

```
print("\nUnique values in 'Sex' column after formatting:")
```

```
print(data['Sex'].unique())
```

5. Detect and Handle Outliers

Boxplot for the 'Fare' column

```
sns.boxplot(data['Fare'], color='skyblue')
```

```
plt.title('Boxplot of Fare')
```

```
plt.show()
```

Detect outliers using the IQR method

```
Q1 = data['Fare'].quantile(0.25)
```

```
Q3 = data['Fare'].quantile(0.75)
```

```
IQR = Q3 - Q1
```

```
lower_bound = Q1 - 1.5 * IQR
```

```
upper_bound = Q3 + 1.5 * IQR
```

Capping outliers

```
data['Fare'] = np.where(data['Fare'] > upper_bound, upper_bound, np.where(data['Fare'] < lower_bound, lower_bound, data['Fare']))
```

Verify with an updated boxplot

```
sns.boxplot(data['Fare'], color='lightgreen')
```

```
plt.title('Boxplot of Fare (After Handling Outliers)')
```



```
plt.show()
```

6. Save the Cleaned Dataset

Save the cleaned dataset

```
data.to_csv('cleaned_titanic.csv', index=False)
```

```
print("\nCleaned dataset saved as 'cleaned_titanic.csv'") .
```



1b. Load a dataset, calculate descriptive summary statistics, create visualizations using different graphs, and identify potential features and target variables Note: Explore Univariate and Bivariate graphs (Matplotlib) and Seaborn for visualization

1. Import Necessary Libraries

Import required libraries

```
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
import seaborn as sns
```

2. Load the Dataset

Load the dataset from the URL

```
url = "https://raw.githubusercontent.com/mwaskom/seaborn-data/master/iris.csv" data =  
pd.read_csv(url)
```

Display the first few rows

```
print("First 5 rows of the dataset:")  
print(data.head())
```

3. Calculate Descriptive Summary Statistics

Dataset information

```
print("\nDataset Info:")  
print(data.info())
```

Summary statistics for numerical columns

```
print("\nDescriptive Statistics for Numerical Columns:")  
print(data.describe())
```

Check unique values for categorical columns

```
print("\nUnique values in 'species' column:")  
print(data['species'].value_counts())
```

4. Univariate Analysis

Histograms for numerical columns

```
data.hist(figsize=(10, 8), color='skyblue', edgecolor='black')  
  
plt.suptitle("Histograms of Numerical Features")  
  
plt.show()
```

Bar plot for 'species' column

```
sns.countplot(x='species', data=data, palette='pastel')  
  
plt.title("Count of Each Species")  
  
plt.show()
```

5. Bivariate Analysis

Scatter plot for two features

```
plt.figure(figsize=(8, 6))  
  
plt.scatter(data['sepal_length'], data['sepal_width'], alpha=0.7, c='blue') plt.title("Sepal  
Length vs Sepal Width")  
  
plt.xlabel("Sepal Length")  
plt.ylabel("Sepal Width")  
  
plt.show()
```

Pairplot to visualize relationships between features

```
sns.pairplot(data, hue='species', palette='husl', diag_kind='kde')  
  
plt.suptitle("Pairplot of Features by Species", y=1.02)  
  
plt.show()
```

Boxplot for petal_length across species

```
sns.boxplot(x='species', y='petal_length', data=data, palette='Set3')  
  
plt.title("Boxplot of Petal Length by Species")  
  
plt.show()
```

6. Identify Potential Features and Target Variables

Separate features and target

```
features = data.drop(columns=['species']) # Drop the target column
```

```
target = data['species'] # Target variable
```

```
print("\nFeatures:")
```

```
print(features.head())
```

```
print("\nTarget:")
```

```
print(target.head())
```

Visualize target distribution

```
sns.countplot(x=target, palette='viridis')
```

```
plt.title("Target Variable Distribution")
```





```
plt.show()
```



7. Save the Cleaned and Processed Dataset


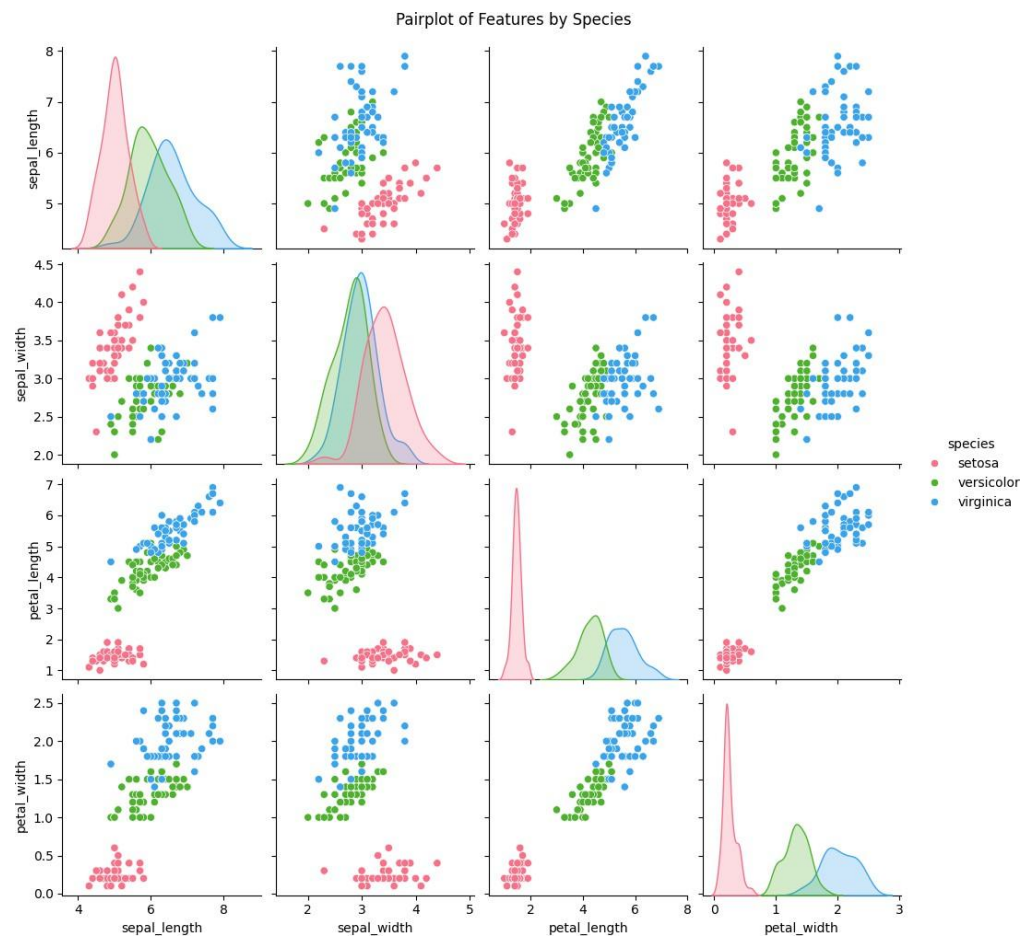
Save the dataset

```
data.to_csv('processed_iris.csv', index=False) print("\nProcessed dataset saved as  
'processed_iris.csv'")
```

Histograms of Numerical Features







1c. Create or Explore datasets to use all pre-processing routines like label encoding, scaling, and binarization.

1. Import Necessary Libraries

Import required libraries

```
import pandas as pd
```

```
import numpy as np
```

```
from sklearn.preprocessing import LabelEncoder, MinMaxScaler, StandardScaler, Binarizer
```

2. Create or Load a Dataset

Create a sample dataset

```
data = pd.DataFrame({  
    'Category': ['A', 'B', 'C', 'A', 'B', 'C'], # Categorical variable  
    'Age': [23, 45, 31, 22, 35, 30], # Numerical variable  
    'Income': [50000, 60000, 70000, 80000, 90000, 100000], # Numerical variable  
    'Has_Car': ['Yes', 'No', 'Yes', 'No', 'Yes', 'No'] # Binary categorical variable })
```

Display the dataset

```
print("Sample Dataset:")
```

```
print(data)
```

3. Apply Pre-Processing Routines

Label Encoding for 'Category' column

```
label_encoder = LabelEncoder()
```

```
data['Category_Encoded'] = label_encoder.fit_transform(data['Category'])
```

Label Encoding for binary column 'Has_Car'

```
data['Has_Car_Encoded'] = label_encoder.fit_transform(data['Has_Car'])
```

```
print("\nAfter Label Encoding:")
```

```
print(data)
```

Min-Max Scaling for 'Income'

```
min_max_scaler = MinMaxScaler()
```

```
data['Income_MinMax'] = min_max_scaler.fit_transform(data[['Income']]) # Standard Scaling  
for 'Age'
```

```

standard_scaler = StandardScaler()

data['Age_Standardized'] = standard_scaler.fit_transform(data[['Age']])

print("\nAfter Scaling:")

print(data)

# Binarization for 'Income' with a threshold of 75,000

binarizer = Binarizer(threshold=75000)

data['Income_Binary'] = binarizer.fit_transform(data[['Income']])

print("\nAfter Binarization:")

print(data)

```

4. Save the Processed Dataset

Save the processed dataset

```

data.to_csv('processed_data.csv', index=False)

print("\nProcessed dataset saved as 'processed_data.csv'")

```

⇒ Sample Dataset:

	Category	Age	Income	Has_Car
0	A	23	50000	Yes
1	B	45	60000	No
2	C	31	70000	Yes
3	A	22	80000	No
4	B	35	90000	Yes
5	C	30	100000	No

⇒ After Label Encoding:

	Category	Age	Income	Has_Car	Category_Encoded	Has_Car_Encoded
0	A	23	50000	Yes	0	1
1	B	45	60000	No	1	0
2	C	31	70000	Yes	2	1
3	A	22	80000	No	0	0
4	B	35	90000	Yes	1	1
5	C	30	100000	No	2	0



After Scaling:

	Category	Age	Income	Has_Car	Category_Encoded	Has_Car_Encoded	\
0	A	23	50000	Yes	0	1	
1	B	45	60000	No	1	0	
2	C	31	70000	Yes	2	1	
3	A	22	80000	No	0	0	
4	B	35	90000	Yes	1	1	
5	C	30	100000	No	2	0	

	Income_MinMax	Age_Standardized
0	0.0	-1.035676
1	0.2	1.812434
2	0.4	0.000000
3	0.6	-1.165136
4	0.8	0.517838
5	1.0	-0.129460



After Binarization:

	Category	Age	Income	Has_Car	Category_Encoded	Has_Car_Encoded	\
0	A	23	50000	Yes	0	1	
1	B	45	60000	No	1	0	
2	C	31	70000	Yes	2	1	
3	A	22	80000	No	0	0	
4	B	35	90000	Yes	1	1	
5	C	30	100000	No	2	0	

	Income_MinMax	Age_Standardized	Income_Binary
0	0.0	-1.035676	0
1	0.2	1.812434	0
2	0.4	0.000000	0
3	0.6	-1.165136	1
4	0.8	0.517838	1
5	1.0	-0.129460	1

2: Testing Hypothesis

AIM: Implement and demonstrate the FIND-S algorithm for finding the most specific hypothesis based on a given set of training data samples. Read the training data from a. CSV file and generate the final specific hypothesis. (Create your dataset)

1. Import Necessary Libraries

Import required libraries

```
import pandas as pd
```

```
import numpy as np
```

2. Create the Dataset and Save it as CSV

Create a synthetic dataset

```
data = {  
    'Sky': ['Sunny', 'Sunny', 'Rainy', 'Sunny', 'Rainy'],  
    'Temperature': ['Warm', 'Cold', 'Warm', 'Warm', 'Cold'],  
    'Humidity': ['Normal', 'High', 'High', 'Normal', 'Normal'],  
    'Wind': ['Strong', 'Strong', 'Weak', 'Strong', 'Weak'],  
    'Water': ['Warm', 'Warm', 'Cool', 'Warm', 'Cool'],  
    'Forecast': ['Same', 'Same', 'Change', 'Same', 'Change'],  
    'Condition': ['Yes', 'No', 'No', 'Yes', 'No'] # Target variable  
}
```

Convert the dataset to a DataFrame

```
df = pd.DataFrame(data)
```

Save the dataset to a CSV file

```
df.to_csv('training_data.csv', index=False)
```

Display the dataset

```
print("Dataset:")
```

```
print(df)
```

3. Load the Dataset

Load the dataset from CSV

```
dataset = pd.read_csv('training_data.csv')
```

Display the dataset

```
print("\nLoaded Dataset:")  
print(dataset)
```

4. Define the FIND-S Algorithm

```
def find_s(training_data):
```

```
    # Extract the features and target
```

```
    features = training_data.iloc[:, :-1].values # All columns except the last
```

```
    target = training_data.iloc[:, -1].values # Last column (target variable)
```

```
    # Initialize the most specific hypothesis
```

```
    hypothesis = ['∅'] * features.shape[1]
```

```
    # Iterate through each example in the dataset
```

```
    for i, example in enumerate(features):
```

```
        if target[i] == 'Yes': # Consider only positive examples for j in range(len(hypothesis)):
```

```
            if hypothesis[j] == '∅': # Update the hypothesis initially hypothesis[j] = example[j]
```

```
            elif hypothesis[j] != example[j]: # Generalize if inconsistent hypothesis[j] = '?'
```

```
    return hypothesis
```

5. Run the FIND-S Algorithm

```
# Apply the FIND-S algorithm
```

```
final_hypothesis = find_s(dataset)
```

```
# Display the final specific hypothesis
```

```
print("\nFinal Specific Hypothesis:")
```

```
print(final_hypothesis)
```



Dataset:

	Sky	Temperature	Humidity	Wind	Water	Forecast	Condition
0	Sunny	Warm	Normal	Strong	Warm	Same	Yes
1	Sunny	Cold	High	Strong	Warm	Same	No
2	Rainy	Warm	High	Weak	Cool	Change	No
3	Sunny	Warm	Normal	Strong	Warm	Same	Yes
4	Rainy	Cold	Normal	Weak	Cool	Change	No



Loaded Dataset:

	Sky	Temperature	Humidity	Wind	Water	Forecast	Condition
0	Sunny	Warm	Normal	Strong	Warm	Same	Yes
1	Sunny	Cold	High	Strong	Warm	Same	No
2	Rainy	Warm	High	Weak	Cool	Change	No
3	Sunny	Warm	Normal	Strong	Warm	Same	Yes
4	Rainy	Cold	Normal	Weak	Cool	Change	No



Final Specific Hypothesis:

['Sunny', 'Warm', 'Normal', 'Strong', 'Warm', 'Same']

3. Linear Models

3a. Simple Linear Regression

Fit a linear regression model on a dataset. Interpret coefficients, make predictions, and evaluate performance using metrics like R-squared and MSE

Step 1: Import Libraries

Import required libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
```

Step 2: Create a Dataset and Save as CSV

Create a sample dataset

```
data = {
    'House_Size': [750, 800, 850, 900, 1000, 1100, 1200, 1300, 1400, 1500],
    'Price': [150000, 160000, 165000, 170000, 180000, 190000, 200000, 210000, 220000,
230000]
}
```

Convert the dataset into a DataFrame

```
df = pd.DataFrame(data)
```

Save to CSV file

```
df.to_csv('house_prices.csv', index=False)
```

Display the dataset

```
print("Dataset:")
```

```
print(df)
```

Step 3: Load the Dataset

Load the dataset

```
dataset = pd.read_csv('house_prices.csv')
```

Display the first few rows

```
print("\nLoaded Dataset:")
```

```
print(dataset.head())
```

Step 4: Split the Dataset into Training and Test Sets

Features and target variable

```
X = dataset[['House_Size']] # Feature: House size
```

```
y = dataset['Price']      # Target: Price
```

Split data into training and testing sets (80% train, 20% test)

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

```
print("\nTraining and Testing Data Sizes:")
```

```
print("Training Data Size:", X_train.shape[0])
```

```
print("Testing Data Size:", X_test.shape[0])
```

Step 5: Fit a Linear Regression Model

Initialize and fit the linear regression model

```
model = LinearRegression()
```

```
model.fit(X_train, y_train)
```

Display the coefficients

```
print("\nModel Coefficients:")
```

```
print("Slope (m):", model.coef_[0])
```

```
print("Intercept (b):", model.intercept_)
```

Step 6: Make Predictions

Predict on the test set

```
y_pred = model.predict(X_test)
```

Display predictions

```
print("\nPredictions on Test Data:")  
print("Actual Prices:", y_test.values)  
print("Predicted Prices:", y_pred)
```

Step 7: Evaluate the Model

Calculate evaluation metrics

```
mse = mean_squared_error(y_test, y_pred)  
r2 = r2_score(y_test, y_pred)
```

Display metrics

```
print("\nModel Performance Metrics:")  
print("Mean Squared Error (MSE):", mse)  
print("R-squared ( $R^2$ ):", r2)
```

Step 8: Visualize the Results

Scatter plot of the training data


```
plt.scatter(X_train, y_train, color='blue', label='Training Data')
```

Plot the regression line

```
plt.plot(X_train, model.predict(X_train), color='red', label='Regression Line')
```

Scatter plot of the test data

```
plt.scatter(X_test, y_test, color='green', label='Test Data')  
plt.title("Simple Linear Regression")  
plt.xlabel("House Size (sq ft)")  
plt.ylabel("Price ($)")  
plt.legend()  
plt.show()
```



3b. Multiple Linear Regression

Extend linear regression to multiple feature. Handle feature selection and potential multicollinearity

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
from statsmodels.stats.outliers_influence import variance_inflation_factor
from sklearn.preprocessing import LabelEncoder # Import LabelEncoder
from sklearn.impute import SimpleImputer
```

Load dataset

```
from google.colab import files
uploaded = files.upload() # Upload your CSV file
```

Read the CSV file

```
data = pd.read_csv(list(uploaded.keys())[0])
```

Display the first few rows

```
print(data.head())
```

Check for null values and basic statistics

```
print(data.info())
print(data.describe())
```

Define a function to calculate VIF

```
def calculate_vif(df):
```

Select only numeric features for VIF calculation

```
    numeric_df = df.select_dtypes(include=np.number)
```

Drop rows with infinite or missing values

```
numeric_df = numeric_df.replace([np.inf, -np.inf], np.nan).dropna()

vif_data = pd.DataFrame()

vif_data["feature"] = numeric_df.columns

vif_data["VIF"] = [variance_inflation_factor(numeric_df.values, i) for i in
range(numeric_df.shape[1])]

return vif_data
```

Selecting features and target variable

```
X = data.drop("Survived", axis=1) # Changed 'y' to 'Survived'

y = data["Survived"]
```

Handle categorical features (e.g., using Label Encoding)

```
for col in X.select_dtypes(include=['object']).columns:
```

```
    le = LabelEncoder()

    X[col] = le.fit_transform(X[col])
```

Impute missing values using the mean (you can choose other strategies)

```
imputer = SimpleImputer(strategy='mean') # Create an imputer instance

X = pd.DataFrame(imputer.fit_transform(X), columns=X.columns) # Impute and update X
```

Calculate VIF for initial features

```
print("VIF before handling multicollinearity:")

print(calculate_vif(X)) # Call the modified function
```

Drop features based on VIF analysis (example: drop 'X1' if VIF is high)

Check if the column exists before dropping

```
if 'X1' in X.columns:

    X = X.drop("X1", axis=1) # Replace 'X1' with the actual high VIF feature name

else:

    print("Column 'X1' not found in the DataFrame.")
```

Recalculate VIF

```
print("VIF after handling multicollinearity:")

print(calculate_vif(X))

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

Initialize and fit the model =

```
LinearRegression()
```

```
model.fit(X_train, y_train)
```

Get coefficients and intercept

```
print("Coefficients:", model.coef_)
```

```
print("Intercept:", model.intercept_)
```

Predictions

```
y_pred = model.predict(X_test)
```

Evaluation metrics

```
rmse = np.sqrt(mean_squared_error(y_test, y_pred))
```

```
r2 = r2_score(y_test, y_pred)
```

```
print(f"RMSE: {rmse}")
```

```
print(f"R^2: {r2}")
```

```
from sklearn.feature_selection import RFE
```

Recursive Feature Elimination

```
rfe = RFE(estimator=LinearRegression(), n_features_to_select=5) # Adjust features
```

```
rfe.fit(X_train, y_train)
```

Selected features

```
print("Selected Features:", X.columns[rfe.support_])
```

Scatter plot of actual vs predicted values

```
plt.scatter(y_test, y_pred)
```

```
plt.xlabel("Actual")
```

```
plt.ylabel("Predicted")
```

```
plt.title("Actual vs Predicted")
```



```
plt.show()
```

Residuals

```
residuals = y_test - y_pred
```

```
sns.histplot(residuals, kde=True)
```

```
plt.title("Residuals Distribution") plt.show()
```



3c. Regularized Linear Models

Implement Regression variants like LASSO and Ridge on any generated dataset

1. Set Up the Environment

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.linear_model import Ridge, Lasso, ElasticNet
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.datasets import make_regression

# Set random seed for reproducibility
np.random.seed(42)
```

2. Generate a Synthetic Dataset

Generate synthetic data

```
X, y = make_regression(
    n_samples=1000, # Number of samples
    n_features=10, # Number of features
    noise=15,      # Add some noise
    random_state=42
)
```

Convert to DataFrame for exploration

```
data = pd.DataFrame(X, columns=[f"X{i}" for i in range(1, 11)])
data["y"] = y
```

Display the first few rows

```
print(data.head())
```

3. Split the Dataset

Split data into training and testing sets

```
X_train, X_test, y_train, y_test = train_test_split(  
    data.drop("y", axis=1), # Features  
    data["y"],             # Target variable  
    test_size=0.2,         # 20% for testing  
    random_state=42  
)
```

4. Train and Evaluate Ridge Regression

Initialize Ridge Regression with a regularization parameter (alpha)

```
ridge = Ridge(alpha=1.0)
```

Train the model

```
ridge.fit(X_train, y_train)
```

Predictions

```
ridge_pred = ridge.predict(X_test)
```

Evaluate Ridge Regression

```
ridge_rmse = np.sqrt(mean_squared_error(y_test, ridge_pred))
```

```
ridge_r2 = r2_score(y_test, ridge_pred)
```

```
print(f"Ridge RMSE: {ridge_rmse}")
```

```
print(f"Ridge R^2: {ridge_r2}")
```

5. Train and Evaluate Lasso Regression

Initialize Lasso Regression

```
lasso = Lasso(alpha=0.1)
```

Train the model

```
lasso.fit(X_train, y_train)
```

Predictions

```
lasso_pred = lasso.predict(X_test)
```

Evaluate Lasso Regression

```
lasso_rmse = np.sqrt(mean_squared_error(y_test, lasso_pred))
lasso_r2 = r2_score(y_test, lasso_pred)
print(f"Lasso RMSE: {lasso_rmse}")
print(f"Lasso R^2: {lasso_r2}")

# Features shrunk to zero
print("Lasso Coefficients:", lasso.coef_)
```

6. Train and Evaluate ElasticNet Regression

Initialize ElasticNet

```
elastic_net = ElasticNet(alpha=0.1, l1_ratio=0.5) # l1_ratio balances L1 and L2 penalties
```

Train the model

```
elastic_net.fit(X_train, y_train)
```

Predictions

```
elastic_net_pred = elastic_net.predict(X_test)
```

Evaluate ElasticNet Regression

```
elastic_net_rmse = np.sqrt(mean_squared_error(y_test, elastic_net_pred))
elastic_net_r2 = r2_score(y_test, elastic_net_pred)
print(f"ElasticNet RMSE: {elastic_net_rmse}")
print(f"ElasticNet R^2: {elastic_net_r2}")
```

7. Compare Results

Collect metrics

```
metrics = pd.DataFrame({
    "Model": ["Ridge", "Lasso", "ElasticNet"],
    "RMSE": [ridge_rmse, lasso_rmse, elastic_net_rmse],
    "R^2": [ridge_r2, lasso_r2, elastic_net_r2]
})


print(metrics)
```

Plot RMSE comparison

```
sns.barplot(data=metrics, x="Model", y="RMSE")
```

```
plt.title("Model RMSE Comparison")
```

```
plt.show()
```



4. Discriminative Models

4a. Logistic Regression : Perform binary classification using logistic regression. Calculate accuracy, precision, recall, and understand the ROC curve."

Step 1: Import Required Libraries

```
# Import necessary libraries
```

```
import numpy as np
```

```
import pandas as pd
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.linear_model import LogisticRegression
```

```
from sklearn.metrics import accuracy_score, precision_score, recall_score, roc_curve, auc
```

```
import matplotlib.pyplot as plt
```

Step 2: Prepare the Dataset

```
from sklearn.datasets import make_classification
```

Create a synthetic dataset

```
X, y = make_classification(n_samples=1000, n_features=10, n_classes=2, random_state=42)
```

Split data into training and testing sets

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

Step 3: Train the Logistic Regression Model

Initialize the logistic regression model

```
logreg = LogisticRegression()
```

Train the model on the training data

```
logreg.fit(X_train, y_train)
```

Step 4: Make Predictions

Predict labels for the test set

```
y_pred = logreg.predict(X_test)
```

Predict probabilities for the ROC curve

```
y_prob = logreg.predict_proba(X_test)[:, 1]
```

Step 5: Evaluate the Model

Calculate metrics

```
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
print(f"Accuracy: {accuracy:.2f}")
print(f"Precision: {precision:.2f}")
print(f"Recall: {recall:.2f}")
```


Step 6: Plot the ROC Curve

Compute ROC curve and AUC

```
fpr, tpr, _ = roc_curve(y_test, y_prob)
roc_auc = auc(fpr, tpr)
```

Plot the ROC curve

```
plt.figure()
plt.plot(fpr, tpr, color='darkorange', lw=2, label=f"ROC curve (AUC = {roc_auc:.2f})")
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic')
plt.legend(loc="lower right")
plt.show()
```



4b .Implement and demonstrate k-nearest Neighbor algorithm. Read the training data from a .CSV file and build the model to classify a test sample. Print both correct and wrong predictions.

Step 1: Import Required Libraries

Import necessary libraries

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import accuracy_score
from google.colab import files
```

Step 2: Create or Upload the CSV File

Check if the user wants to create a dataset or upload one

```
print("Do you have a CSV file to upload? (yes/no)")
response = input().lower()
if response == "yes":
    uploaded = files.upload()
    filename = list(uploaded.keys())[0]
else:
    # Create a synthetic dataset
    from sklearn.datasets import make_classification
    # Generate synthetic data
    X, y = make_classification(
        n_samples=200, n_features=5, n_classes=2, random_state=42
    )
    # Combine features and target into a single DataFrame
    data = pd.DataFrame(X, columns=[f"Feature_{i}" for i in range(X.shape[1])])
    data['Target'] = y
```

Save the dataset to a CSV file

```
filename = "synthetic_data.csv"
data.to_csv(filename, index=False)
print(f"Synthetic dataset saved as {filename}.")
```

Step 3: Load the CSV File into a DataFrame

Load the dataset into a DataFrame

```
data = pd.read_csv(filename)
```

Display the first few rows of the dataset

```
print("Loaded Dataset:")
print(data.head())
```

Step 4: Preprocess the Data

Separate features (X) and labels (y)

```
X = data.iloc[:, :-1].values # All columns except the last one
y = data.iloc[:, -1].values # Last column as the target
```

Split the dataset into training and testing sets (80% train, 20% test)

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

Step 5: Train the k-NN Model

Initialize the k-NN model with k=3

```
knn = KNeighborsClassifier(n_neighbors=3)
```

Train the model on the training data

```
knn.fit(X_train, y_train)
```

Step 6: Predict Test Samples

Predict the labels for the test set

```
y_pred = knn.predict(X_test)
```


Step 7: Evaluate and Print Predictions

Calculate and display the accuracy

```
accuracy = accuracy_score(y_test, y_pred)
print(f"\nModel Accuracy: {accuracy:.2f}\n")
```

Display correct and incorrect predictions

```
print("Correct Predictions:")
for i in range(len(y_test)):
    if y_pred[i] == y_test[i]:
        print(f"Sample {i}: Predicted={y_pred[i]}, Actual={y_test[i]}")
print("\nIncorrect Predictions:")
for i in range(len(y_test)):
    if y_pred[i] != y_test[i]:
        print(f"Sample {i}: Predicted={y_pred[i]}, Actual={y_test[i]}")
```

Output :

```
[ ]
Model Accuracy: 0.88

Correct Predictions:
Sample 0: Predicted=0, Actual=0
Sample 1: Predicted=1, Actual=1
Sample 2: Predicted=1, Actual=1
Sample 3: Predicted=0, Actual=0
Sample 4: Predicted=1, Actual=1
Sample 5: Predicted=1, Actual=1
Sample 6: Predicted=0, Actual=0
Sample 7: Predicted=0, Actual=0
Sample 9: Predicted=1, Actual=1
Sample 10: Predicted=1, Actual=1
Sample 11: Predicted=1, Actual=1
Sample 12: Predicted=0, Actual=0
Sample 13: Predicted=0, Actual=0
Sample 14: Predicted=0, Actual=0
Sample 15: Predicted=0, Actual=0
Sample 16: Predicted=0, Actual=0
Sample 17: Predicted=1, Actual=1
Sample 18: Predicted=1, Actual=1
Sample 19: Predicted=0, Actual=0
Sample 20: Predicted=0, Actual=0
Sample 22: Predicted=1, Actual=1
Sample 23: Predicted=1, Actual=1
Sample 24: Predicted=1, Actual=1
Sample 25: Predicted=1, Actual=1
Sample 26: Predicted=1, Actual=1
Sample 27: Predicted=0, Actual=0
Sample 28: Predicted=0, Actual=0
Sample 30: Predicted=1, Actual=1
Sample 31: Predicted=1, Actual=1
Sample 32: Predicted=1, Actual=1
Sample 34: Predicted=0, Actual=0
Sample 35: Predicted=1, Actual=1
Sample 36: Predicted=1, Actual=1
Sample 38: Predicted=1, Actual=1
Sample 39: Predicted=1, Actual=1
```

Incorrect Predictions:

Sample 8: Predicted=1, Actual=0

Sample 21: Predicted=1, Actual=0

Sample 29: Predicted=0, Actual=1

Sample 33: Predicted=1, Actual=0

Sample 37: Predicted=1, Actual=0

4c. Build a decision tree classifier or regressor. Control hyperparameters like tree depth to avoid overfitting. Visualize the tree.

Step 1: Import Required Libraries

Import necessary libraries

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier, DecisionTreeRegressor, plot_tree
from sklearn.metrics import accuracy_score, mean_squared_error
import matplotlib.pyplot as plt
from google.colab import files
```

Step 2: Create or Upload the CSV File

Check if the user wants to upload a file or generate one

```
print("Do you have a CSV file to upload? (yes/no)")
response = input().lower()
if response == "yes":
    # Upload the CSV file
    uploaded = files.upload()
    filename = list(uploaded.keys())[0]
else:
    # Generate synthetic data (classification or regression)
    from sklearn.datasets import make_classification, make_regression
    print("Choose a task: (1) Classification (2) Regression")
    task = int(input())
    if task == 1:
        # Generate synthetic classification data
        X, y = make_classification(n_samples=200, n_features=5, random_state=42)
        task_type = "classification"
```

else:

Generate synthetic regression data

X, y = make_regression(n_samples=200, n_features=5, random_state=42)

task_type = "regression"

Combine features and target into a single DataFrame

data = pd.DataFrame(X, columns=[f"Feature_{i}" for i in range(X.shape[1])])

data['Target'] = y

Save the dataset to a CSV file

filename = "synthetic_data.csv"

data.to_csv(filename, index=False)

print(f"Synthetic {task_type} dataset saved as {filename}.")

Step 3: Load the Dataset

Load the dataset

data = pd.read_csv(filename)

Display the first few rows of the dataset

print("Dataset Preview:")

print(data.head())

Step 4: Preprocess the Data

Separate features and target

X = data.iloc[:, :-1].values # All columns except the last one

y = data.iloc[:, -1].values # Last column as the target

Split data into training and testing sets

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

Step 5: Build the Decision Tree

Define the tree depth to avoid overfitting

max_depth = 3

Initialize the model

```
if task_type == "classification":  
    model = DecisionTreeClassifier(max_depth=max_depth, random_state=42)  
else:  
    model = DecisionTreeRegressor(max_depth=max_depth, random_state=42)
```

Train the model

```
model.fit(X_train, y_train)
```

Step 6: Make Predictions

Predict on the test set

```
y_pred = model.predict(X_test)
```

Evaluate the model

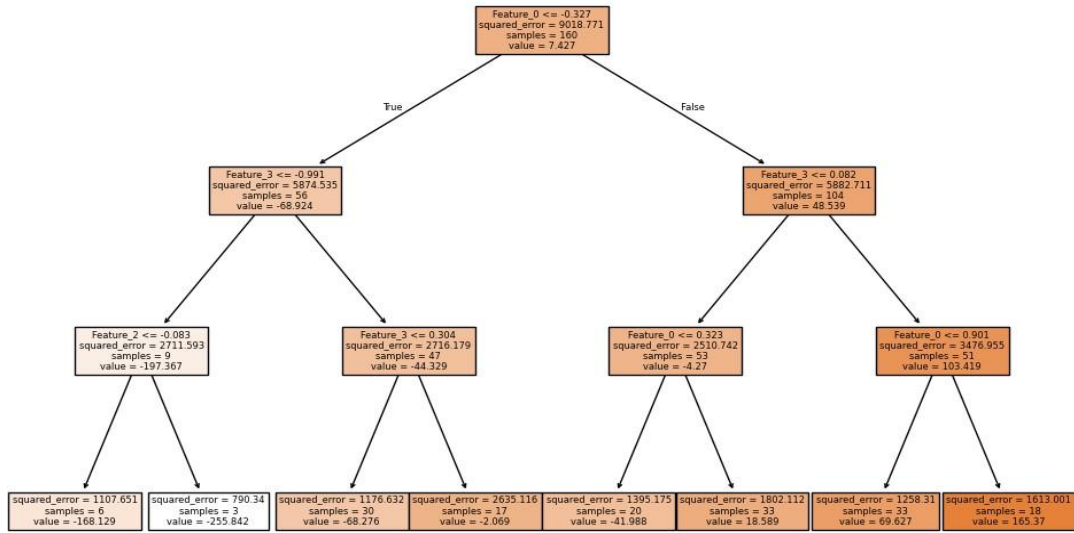
```
if task_type == "classification":  
    accuracy = accuracy_score(y_test, y_pred)  
    print(f"Accuracy: {accuracy:.2f}")  
else:  
    mse = mean_squared_error(y_test, y_pred)  
    print(f"Mean Squared Error: {mse:.2f}")
```

Step 7: Visualize the Tree

Visualize the decision tree

```
plt.figure(figsize=(12, 8))  
  
plot_tree(model, feature_names=data.columns[:-1], class_names=str(np.unique(y)) if  
task_type == "classification" else None, filled=True)  
  
plt.title("Decision Tree Visualization")  
  
plt.show()
```

Decision Tree Visualization



4d. Implement a Support Vector Machine for any relevant dataset.

Step 1: Import Required Libraries

Import necessary libraries

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, classification_report
from google.colab import files
```

Step 2: Create or Upload a Dataset

Check if the user wants to upload a file or generate one

```
print("Do you have a CSV file to upload? (yes/no)")
```

```
response = input().lower()
```

```
if response == "yes":
```

Upload the CSV file

```
    uploaded = files.upload()
```

```
    filename = list(uploaded.keys())[0]
```

```
else:
```

Generate synthetic classification data

```
from sklearn.datasets import make_classification
```

```
X, y = make_classification(n_samples=200, n_features=5, n_classes=2, random_state=42)
```

Combine features and target into a DataFrame

```
data = pd.DataFrame(X, columns=[f"Feature_{i}" for i in range(X.shape[1])])
```

```
data['Target'] = y
```

Save the synthetic dataset to a CSV file

```
filename = "synthetic_data.csv"
```

```
data.to_csv(filename, index=False)

print(f"Synthetic dataset saved as {filename}.")
```

Step 3: Load the Dataset

Load the dataset into a DataFrame

```
data = pd.read_csv(filename)
```

Display the first few rows of the dataset

```
print("Dataset Preview:")

print(data.head())
```

Step 4: Preprocess the Data

Separate features (X) and target (y)

```
X = data.iloc[:, :-1].values # All columns except the last one
y = data.iloc[:, -1].values # Last column as the target
```

Split the dataset into training (80%) and testing (20%) sets

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

Step 5: Train the Support Vector Machine

Initialize the SVM model (use RBF kernel as default)

```
svm_model = SVC(kernel='rbf', C=1.0, gamma='scale', random_state=42)
```

Train the SVM model on the training data

```
svm_model.fit(X_train, y_train)
```

Step 6: Make Predictions

Predict the labels for the test set

```
y_pred = svm_model.predict(X_test)
```

Step 7: Evaluate the Model

Calculate and print the accuracy


```
accuracy = accuracy_score(y_test, y_pred)
print(f"Model Accuracy: {accuracy:.2f}")
```

Print a detailed classification report

```
print("\nClassification Report:")
print(classification_report(y_test, y_pred))
```

Step 8: Visualize the Decision Boundary (Optional for 2D Data)

```
import matplotlib.pyplot as plt
```

Generate 2D synthetic data

```
from sklearn.datasets import make_blobs
```

```
X, y = make_blobs(n_samples=100, centers=2, random_state=42, cluster_std=1.5)
```

Fit the SVM on this data

```
svm_model.fit(X, y)
```

Plot the decision boundary

```
plt.figure(figsize=(8, 6))
```

```
plt.scatter(X[:, 0], X[:, 1], c=y, cmap='coolwarm', edgecolor='k')
```

Create a grid to evaluate the model

```
xx, yy = np.meshgrid(np.linspace(X[:, 0].min(), X[:, 0].max(), 100),
                     np.linspace(X[:, 1].min(), X[:, 1].max(), 100))
```

```
Z = svm_model.decision_function(np.c_[xx.ravel(), yy.ravel()])
```

```
Z = Z.reshape(xx.shape)
```

Plot the decision boundary and margins


```
plt.contour(xx, yy, Z, levels=[-1, 0, 1], linestyles=['--', '-', '--'], colors='k')
```

```
plt.title("SVM Decision Boundary")
```

```
plt.xlabel("Feature 1")
```

```
plt.ylabel("Feature 2")
```

```
plt.show()
```



4e. Train a random forest ensemble. Experiment with the number of trees and feature sampling. Compare performance to a single decision tree.

Step 1: Import Required Libraries

Import necessary libraries

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report
from google.colab import files
```

Step 2: Create or Upload a Dataset

Check if the user wants to upload a file or generate one

```
print("Do you have a CSV file to upload? (yes/no)")
response = input().lower()
if response == "yes":
    # Upload the CSV file
    uploaded = files.upload()
    filename = list(uploaded.keys())[0]
else:
```

Generate synthetic classification data

```
from sklearn.datasets import make_classification
X, y = make_classification(n_samples=300, n_features=10, n_classes=2, random_state=42)
```

Combine features and target into a DataFrame

```
data = pd.DataFrame(X, columns=[f"Feature_{i}" for i in range(X.shape[1])])
data['Target'] = y
```

Save the synthetic dataset to a CSV file

```
filename = "synthetic_data.csv"
data.to_csv(filename, index=False)
```

```
print(f"Synthetic dataset saved as {filename}.")
```

Step 3: Load the Dataset

Load the dataset

```
data = pd.read_csv(filename)
```

Display the first few rows of the dataset

```
print("Dataset Preview:")
```

```
print(data.head())
```

Step 4: Preprocess the Data

Separate features (X) and target (y)

```
X = data.iloc[:, :-1].values # All columns except the last one
```

```
y = data.iloc[:, -1].values # Last column as the target
```

Split the dataset into training (80%) and testing (20%) sets

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

Step 5: Train a Single Decision Tree Classifier

Initialize and train the Decision Tree model

```
decision_tree = DecisionTreeClassifier(random_state=42)
```

```
decision_tree.fit(X_train, y_train)
```

Predict and evaluate

```
y_pred_tree = decision_tree.predict(X_test)
```

```
accuracy_tree = accuracy_score(y_test, y_pred_tree)
```

```
print(f"Decision Tree Accuracy: {accuracy_tree:.2f}")
```

Step 6: Train a Random Forest Classifier

Initialize the Random Forest model with hyperparameter tuning

```
random_forest = RandomForestClassifier(n_estimators=100, max_features='sqrt',  
random_state=42)
```

Train the model

```
random_forest.fit(X_train, y_train)
```

Predict and evaluate

```
y_pred_rf = random_forest.predict(X_test)
```

```
accuracy_rf = accuracy_score(y_test, y_pred_rf)
```

```
print(f"Random Forest Accuracy (100 trees, sqrt features): {accuracy_rf:.2f}")
```

Step 7: Experiment with Random Forest Hyperparameters

Experiment with fewer trees and different feature sampling

```
rf_experiment = RandomForestClassifier(n_estimators=50, max_features=3,  
random_state=42)
```

```
rf_experiment.fit(X_train, y_train)
```

Predict and evaluate

```
y_pred_rf_exp = rf_experiment.predict(X_test)
```

```
accuracy_rf_exp = accuracy_score(y_test, y_pred_rf_exp)
```

```
print(f"Random Forest Accuracy (50 trees, max_features=3): {accuracy_rf_exp:.2f}")
```

Step 8: Compare the Models

```
print("\nModel Comparison:")
```

```
print(f"Decision Tree Accuracy: {accuracy_tree:.2f}")
```

```
print(f"Random Forest Accuracy (100 trees): {accuracy_rf:.2f}")
```

```
print(f"Random Forest Accuracy (50 trees, max_features=3): {accuracy_rf_exp:.2f}")
```

Step 9: Visualize Feature Importance (Optional)

```
import matplotlib.pyplot as plt
```

Extract feature importance from the Random Forest model


```
feature_importances = random_forest.feature_importances_
```

Plot the feature importance

```
plt.figure(figsize=(10, 6))
```

```
plt.bar(range(len(feature_importances)), feature_importances, tick_label=data.columns[:-1])
```

```
plt.title("Feature Importance in Random Forest")  
plt.xlabel("Features")  
plt.ylabel("Importance")  
plt.show()
```



4f. Implement a gradient boosting machine (e.g., XGBoost). Tune hyperparameters and explore feature importance.

Step 1: Import Required Libraries

```
# Import necessary libraries

import pandas as pd

import numpy as np

from sklearn.model_selection import train_test_split, GridSearchCV

from sklearn.metrics import accuracy_score, classification_report

from xgboost import XGBClassifier, plot_importance

import matplotlib.pyplot as plt

from google.colab import files
```

Step 2: Create or Upload a Dataset

Check if the user wants to upload a file or generate one

```
print("Do you have a CSV file to upload? (yes/no)")

response = input().lower()

if response == "yes":

    # Upload the CSV file

    uploaded = files.upload()

    filename = list(uploaded.keys())[0]

else:

    # Generate synthetic classification data

    from sklearn.datasets import make_classification

    X, y = make_classification(n_samples=300, n_features=10, n_classes=2, random_state=42)

    # Combine features and target into a DataFrame

    data = pd.DataFrame(X, columns=[f"Feature_{i}" for i in range(X.shape[1])])

    data['Target'] = y

    # Save the synthetic dataset to a CSV file

    filename = "synthetic_data.csv"
```

```
data.to_csv(filename, index=False)

print(f"Synthetic dataset saved as {filename}.")
```

Step 3: Load the Dataset

Load the dataset

```
data = pd.read_csv(filename)
```

Display the first few rows of the dataset

```
print("Dataset Preview:")

print(data.head())
```

Step 4: Preprocess the Data

Separate features (X) and target (y)

```
X = data.iloc[:, :-1].values # All columns except the last one
y = data.iloc[:, -1].values # Last column as the target
```

Split the dataset into training (80%) and testing (20%) sets

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

Step 5: Train a Basic XGBoost Model

Initialize and train the XGBoost model with default parameters

```
xgb = XGBClassifier(random_state=42)

xgb.fit(X_train, y_train)
```

Predict and evaluate the model

```
y_pred = xgb.predict(X_test)

accuracy = accuracy_score(y_test, y_pred)

print(f"XGBoost Accuracy (Default Parameters): {accuracy:.2f}")
```

Step 6: Tune Hyperparameters with GridSearchCV

Define a grid of hyperparameters

```
param_grid = {
```



```

'n_estimators': [50, 100, 150],
'learning_rate': [0.01, 0.1, 0.2],
'max_depth': [3, 5, 7]
}

# Initialize GridSearchCV
grid_search = GridSearchCV(estimator=XGBClassifier(random_state=42),
                           param_grid=param_grid,
                           scoring='accuracy',
                           cv=3,
                           verbose=1)

# Fit the model with grid search
grid_search.fit(X_train, y_train)

# Best parameters from GridSearch
print(f"Best Parameters: {grid_search.best_params_}")

# Train the final model with the best parameters
best_xgb = grid_search.best_estimator_

# Predict and evaluate
y_pred_best = best_xgb.predict(X_test)
accuracy_best = accuracy_score(y_test, y_pred_best)
print(f"XGBoost Accuracy (Tuned Parameters): {accuracy_best:.2f}")

```

Step 7: Explore Feature Importance

Plot feature importance for the tuned model

```

plt.figure(figsize=(10, 6))

plot_importance(best_xgb, importance_type='weight', xlabel="Importance",
                ylabel="Features")

plt.title("XGBoost Feature Importance")

plt.show()

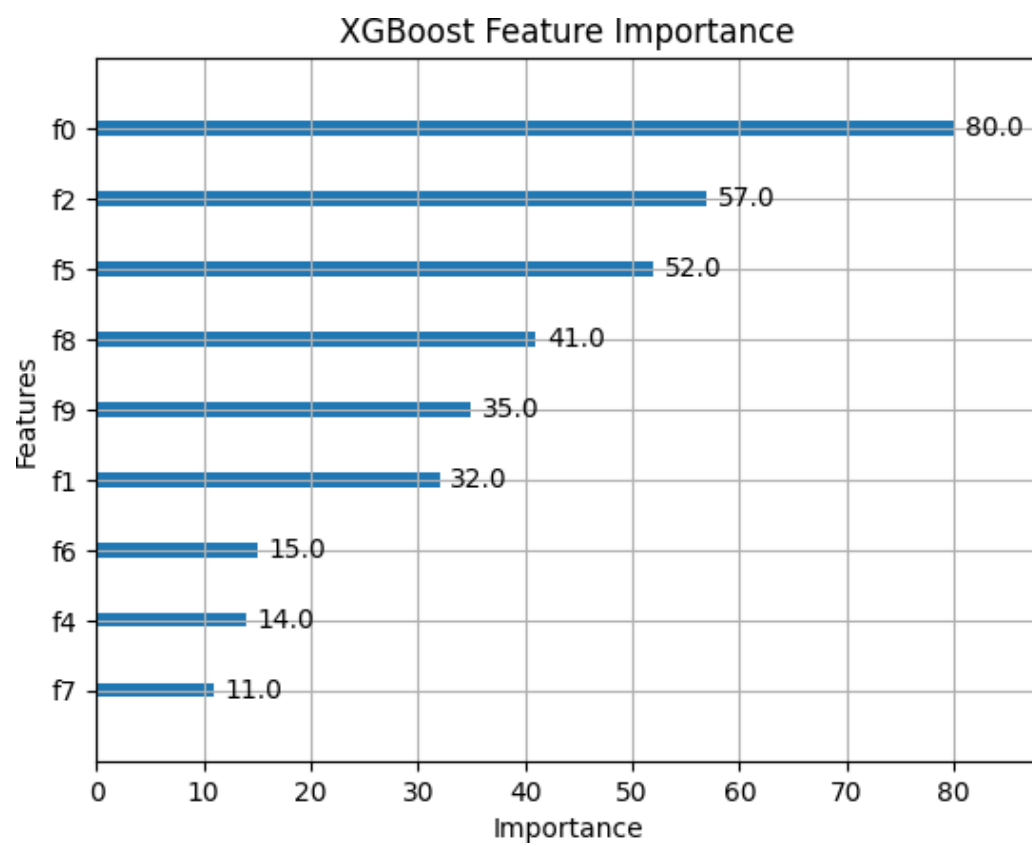
```

Step 8: Evaluate the Model

Print a detailed classification report

```
print("Classification Report:")
```

```
print(classification_report(y_test, y_pred_best))
```



5a. Implement and demonstrate the working of a Naive Bayesian classifier using a sample data set. Build the model to classify a test sample.

Step 1: Import Required Libraries

Import necessary libraries

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, classification_report
from sklearn.naive_bayes import GaussianNB
from google.colab import files
```

Step 2: Create or Upload a Dataset

Ask if the user wants to upload a file or generate one

```
print("Do you have a CSV file to upload? (yes/no)")
response = input().lower()
if response == "yes":
    # Upload the CSV file
    uploaded = files.upload()
    filename = list(uploaded.keys())[0]
else:
    # Generate synthetic classification data
    from sklearn.datasets import make_classification
    X, y = make_classification(n_samples=300, n_features=8, n_classes=2, random_state=42)
    # Combine features and target into a DataFrame
    data = pd.DataFrame(X, columns=[f"Feature_{i}" for i in range(X.shape[1])])
    data['Target'] = y
    # Save the synthetic dataset to a CSV file
    filename = "synthetic_naive_bayes_data.csv"
    data.to_csv(filename, index=False)
```

```
print(f"Synthetic dataset saved as {filename}.")
```

Step 3: Load the Dataset

Load the dataset

```
data = pd.read_csv(filename)
```

Display the first few rows of the dataset

```
print("Dataset Preview:")
```

```
print(data.head())
```

Step 4: Preprocess the Data

Separate features (X) and target (y)

```
X = data.iloc[:, :-1].values # All columns except the last one
```

```
y = data.iloc[:, -1].values # Last column as the target
```

Split the dataset into training (80%) and testing (20%) sets

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

Step 5: Train a Naive Bayes Classifier

Initialize the Gaussian Naive Bayes classifier

```
naive_bayes = GaussianNB()
```

Train the model

```
naive_bayes.fit(X_train, y_train)
```

Step 6: Make Predictions and Evaluate

Predict on the test set

```
y_pred = naive_bayes.predict(X_test)
```

Evaluate the model

```
accuracy = accuracy_score(y_test, y_pred)
```

```
print(f"Naive Bayes Accuracy: {accuracy:.2f}")
```

Detailed classification report

```
print("Classification Report:")  
print(classification_report(y_test, y_pred))
```

Step 7: Test the Model with a Custom Sample

Define a sample test input (replace with meaningful values based on your dataset)

```
test_sample = [X_test[0]] # Taking the first test sample for demonstration
```

Predict the class for the test sample

```
predicted_class = naive_bayes.predict(test_sample)  
print(f"Test Sample: {test_sample}")  
print(f"Predicted Class: {predicted_class[0]}")
```

```
Dataset Preview:
```

	Feature_0	Feature_1	Feature_2	Feature_3	Feature_4	Feature_5	\
0	-1.274158	1.317988	-2.423879	0.906946	-1.583903	-0.331811	
1	1.607963	-1.649959	0.299293	-0.891720	1.301741	1.508502	
2	-0.154167	0.161033	2.210523	0.139400	-0.557492	0.087713	
3	-0.920991	0.949136	-1.613561	0.588410	1.471170	-0.529287	
4	1.013304	-1.038578	-0.305225	-0.539334	-0.609512	1.048078	

	Feature_6	Feature_7	Target
0	-0.452306	0.760415	1
1	0.742095	1.561511	0
2	0.963879	-1.369803	0
3	-1.371901	-0.209324	0
4	-1.065114	-0.186971	0

```
Test Sample: [array([-0.90320608,  0.9220511 , -1.32308979,  0.41081065,  1.64201516,  
                    -1.23559176, -0.63896175,  1.00981709])]  
Predicted Class: 1
```

5b. Implement Hidden Markov Models using hmmlearn

Step 1: Install Required Libraries

Install hmmlearn

```
!pip install hmmlearn
```

Step 2: Import Required Libraries

Import necessary libraries

```
import numpy as np
import pandas as pd
from hmmlearn import hmm
import matplotlib.pyplot as plt
```

Step 3: Create or Load a Dataset

Generate synthetic observable data

```
np.random.seed(42)
```

Create a sequence of observations and hidden states

```
observations = np.random.choice(['A', 'B', 'C'], size=100, p=[0.5, 0.3, 0.2])
hidden_states = np.random.choice(['X', 'Y'], size=100, p=[0.6, 0.4])
```

Save the data in a DataFrame for analysis

```
data = pd.DataFrame({'Observations': observations, 'Hidden States': hidden_states})
print("Generated Data:")
print(data.head())
```

Step 4: Encode Observations

Encode the observations into integers

```
observation_mapping = {obs: idx for idx, obs in enumerate(np.unique(observations))}
encoded_observations = np.array([observation_mapping[obs] for obs in observations])
```

Print the mapping

```
print("Observation Encoding:")  
print(observation_mapping)
```

Step 5: Initialize and Configure the HMM

Initialize the HMM model

```
n_states = 2 # Number of hidden states  
  
n_observations = len(observation_mapping) # Number of unique observations  
  
model = hmm.MultinomialHMM(n_components=n_states, random_state=42, n_iter=100,  
tol=0.01)
```

Define start probabilities (initial distribution of states)

```
start_probs = np.array([0.6, 0.4]) # Assumed probabilities  
model.startprob_ = start_probs
```

Define transition probabilities between states

```
trans_probs = np.array([  
    [0.7, 0.3], # From state X  
    [0.4, 0.6], # From state Y  
)  
model.transmat_ = trans_probs
```

Define emission probabilities (probability of observations given states)

```
emission_probs = np.array([  
    [0.5, 0.4, 0.1], # State X emits A, B, C  
    [0.2, 0.3, 0.5], # State Y emits A, B, C  
)  
model.emissionprob_ = emission_probs
```

Print the configured model parameters

```
print("Start Probabilities:", model.startprob_)  
print("Transition Matrix:", model.transmat_)  
print("Emission Probabilities:", model.emissionprob_)
```

Step 6: Train the Model

Reshape the data for HMM (requires 2D array)

```
encoded_observations = encoded_observations.reshape(-1, 1)
```

Fit the model

```
model.fit(encoded_observations)
```

Predict hidden states for the observations

```
predicted_states = model.predict(encoded_observations)
```

Print the predicted states

```
print("Predicted States:")
```

```
print(predicted_states)
```

Step 7: Visualize the Results

Map predicted states back to their original labels

```
state_mapping = {0: 'X', 1: 'Y'}
```

```
predicted_state_labels = [state_mapping[state] for state in predicted_states]
```

Add predicted states to the DataFrame

```
data['Predicted States'] = predicted_state_labels
```

Display the first few rows with predicted states

```
print("Data with Predicted States:")
```

```
print(data.head())
```

Plot the observations and predicted states

```
plt.figure(figsize=(12, 6))
```

```
plt.plot(data['Observations'], label='Observations', marker='o', linestyle='-', alpha=0.7)
```

```
plt.plot(data['Predicted States'], label='Predicted States', marker='x', linestyle='--', alpha=0.7)
```


```
plt.legend()
```

```
plt.title("Observations and Predicted States")
```

```
plt.xlabel("Time")
```

```
plt.ylabel("Value")
```

```
plt.show()
```

6. Probabilistic Models

6a. Implement Bayesian Linear Regression to explore prior and posterior distribution.

Bayesian Linear Regression is a probabilistic approach to linear regression that incorporates uncertainty in the model parameters. Instead of estimating point values for parameters (as in traditional linear regression), we estimate distributions over the parameters.

Step 1: Install Required Libraries

Install necessary libraries

```
!pip install matplotlib seaborn scikit-learn
```

Step 2: Import Required Libraries

Import necessary libraries

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.linear_model import BayesianRidge
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
from google.colab import files
```

Step 3: Create or Upload a Dataset

Upload a CSV file if you have one

```
print("Do you have a CSV file to upload? (yes/no)")
response = input().lower()
if response == "yes":
    # Upload the CSV file
    uploaded = files.upload()
```

```

filename = list(uploaded.keys())[0]
else:
    # Generate synthetic data for demonstration
    np.random.seed(42)
    X = np.random.rand(100, 1) * 10 # Random data between 0 and 10
    y = 2 * X + 1 + np.random.randn(100, 1) * 2 # y = 2x + 1 with some noise
    # Convert to a DataFrame
    data = pd.DataFrame(np.hstack((X, y)), columns=["X", "y"])
    # Save to CSV for convenience
    filename = "synthetic_data.csv"
    data.to_csv(filename, index=False)
    print(f"Synthetic dataset saved as {filename}.")

```

Step 4: Load and Explore the Data

Load the dataset (for CSV file)

```
data = pd.read_csv(filename)
```

Display first few rows

```
print("Dataset Preview:")
print(data.head())
```

Step 5: Preprocess the Data

Separate features (X) and target (y)

```
X = data["X"].values.reshape(-1, 1) # Feature matrix
y = data["y"].values # Target vector
```

Split the dataset into training (80%) and testing (20%) sets

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

Step 6: Implement Bayesian Linear Regression Model

Initialize the BayesianRidge model (which implements Bayesian Linear Regression)

```
bayesian_regressor = BayesianRidge()

# Fit the model on the training data

bayesian_regressor.fit(X_train, y_train)

# Predict on the test data

y_pred = bayesian_regressor.predict(X_test)
```

Step 7: Visualize the Prior and Posterior Distributions

Plot the prior and posterior distributions of the parameters

```
fig, ax = plt.subplots(1, 2, figsize=(12, 6))

# Plot prior distribution (assuming the model starts with a standard prior)

ax[0].set_title("Prior Distribution (Assumed)")

ax[0].hist(np.random.normal(0, 1, 1000), bins=50, alpha=0.7, color='blue', label="Prior")

ax[0].legend()

# Plot posterior distribution (after model fitting)

ax[1].set_title("Posterior Distribution (After Fitting)")

ax[1].hist(bayesian_regressor.coef_, bins=50, alpha=0.7, color='green', label="Posterior")

ax[1].legend()

plt.show()
```

Step 8: Evaluate the Model Performance

Calculate the Mean Squared Error (MSE)

```
mse = mean_squared_error(y_test, y_pred)

print(f"Mean Squared Error (MSE): {mse:.2f}")
```

Step 9: Visualize the Fit of the Model

Plot the true values and the predicted values

```
plt.figure(figsize=(8, 6))

plt.scatter(X_test, y_test, color="blue", label="True values")

plt.plot(X_test, y_pred, color="red", label="Predicted values", linewidth=2)
```

```
plt.title("Bayesian Linear Regression: True vs Predicted Values")



plt.xlabel("X")

plt.ylabel("y")

plt.legend()

plt.show()
```

Mean Squared Error (MSE): 3.9



6b. Implement Gaussian Mixture Models for density estimation and unsupervised clustering.

Step 1: Install Required Libraries

Install required libraries

```
!pip install matplotlib seaborn scikit-learn
```

Step 2: Import Required Libraries

Import necessary libraries

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.mixture import GaussianMixture
from sklearn.model_selection import train_test_split
from google.colab import files
```

Step 3: Create or Upload a Dataset

Ask if the user has a CSV file to upload

```
print("Do you have a CSV file to upload? (yes/no)")
response = input().lower()
if response == "yes":
    # Upload the CSV file
    uploaded = files.upload()
    filename = list(uploaded.keys())[0]
else:
    # Generate synthetic 2D data with two clusters for demonstration
    np.random.seed(42)
    # Generate data for two Gaussian distributions
    X1 = np.random.normal(loc=0, scale=1, size=(300, 2)) # Cluster 1: mean = 0, std = 1
    X2 = np.random.normal(loc=5, scale=1, size=(300, 2)) # Cluster 2: mean = 5, std = 1
```

Stack the data to create a dataset

```
X = np.vstack([X1, X2])
```

Create DataFrame to simulate the CSV file for consistency

```
data = pd.DataFrame(X, columns=["Feature_1", "Feature_2"])
```

```
filename = "synthetic_data.csv"
```

```
data.to_csv(filename, index=False)
```

```
print(f"Synthetic dataset saved as {filename}.")
```

Step 4: Load and Explore the Dataset

Load the dataset (if CSV file is uploaded)

```
data = pd.read_csv(filename)
```

Display the first few rows

```
print("Dataset Preview:")
```

```
print(data.head())
```

Plot the data to visualize its structure

```
sns.scatterplot(data=data, x="Feature_1", y="Feature_2")
```

```
plt.title("Synthetic Data")
```

```
plt.xlabel("Feature 1")
```

```
plt.ylabel("Feature 2")
```

```
plt.show()
```

Step 5: Fit a Gaussian Mixture Model (GMM)

Define the GMM model

```
n_components = 2 # Number of Gaussian distributions (clusters)
```

```
gmm = GaussianMixture(n_components=n_components, covariance_type='full',  
random_state=42)
```

Fit the GMM model to the data

```
gmm.fit(data)
```

Predict the cluster labels for each data point

```
labels = gmm.predict(data)
```

Add the cluster labels to the dataset for visualization

```
data['Cluster'] = labels
```

Plot the clustered data

```
sns.scatterplot(data=data, x="Feature_1", y="Feature_2", hue="Cluster", palette="viridis",  
marker="o")
```

```
plt.title("Gaussian Mixture Model Clustering")
```

```
plt.xlabel("Feature 1")
```

```
plt.ylabel("Feature 2")
```

```
plt.legend()
```

```
plt.show()
```

Step 6: Visualize the Gaussian Mixture Model (GMM) Components

Extract the means and covariances of the Gaussian components

```
means = gmm.means_
```

```
covariances = gmm.covariances_
```

Plot the GMM components on top of the data

```
plt.figure(figsize=(8, 6))
```

Plot data points

```
sns.scatterplot(data=data, x="Feature_1", y="Feature_2", hue="Cluster", palette="viridis",  
marker="o", s=60, alpha=0.7)
```

Plot the GMM ellipses

```
for mean, covar in zip(means, covariances):
```

```
    # Plot the Gaussian components as ellipses
```

```
    v, w = np.linalg.eigh(covar)
```

```
    v = 2.0 * np.sqrt(2.0) * np.sqrt(v) # Scaling factor for the ellipse
```

```
    u = w[0] / np.linalg.norm(w[0]) # Normalize the eigenvector
```

```
    angle = np.arctan(u[1] / u[0])
```

Create the ellipse

```
    angle = angle * 180.0 / np.pi # Convert to degrees
```



```
ellipse = plt.matplotlib.patches.Ellipse(mean, v[0], v[1], angle=angle, color='red',  
alpha=0.3)
```

```
plt.gca().add_patch(ellipse)
```

```
plt.title("GMM Clustering with Gaussian Components")
```

```
plt.xlabel("Feature 1")
```

```
plt.ylabel("Feature 2")
```

```
plt.legend()
```

```
plt.show()
```

Step 7: Model Evaluation (Optional)

Compute the log-likelihood of the data under the fitted GMM model

```
log_likelihood = gmm.score(data)
```

```
print(f"Log-Likelihood of the data: {log_likelihood:.2f}")
```

Step 8: Predict New Data Points

Example of predicting the cluster for new data points

```
new_data = np.array([[1.5, 2.5], [4.5, 5.5], [7.0, 8.0]])
```



```
new_labels = gmm.predict(new_data)
```

Print the predicted clusters for the new data points

```
print("Predicted Clusters for New Data Points:")
```

```
for i, label in enumerate(new_labels):
```

```
    print(f>Data point {new_data[i]} is in Cluster {label}")
```



7. Model Evaluation and Hyperparameter Tuning

a. Implement cross-validation techniques (k-fold, stratified, etc.) for robust model evaluation

1. Import Necessary Libraries

```
import numpy as np
import pandas as pd
from sklearn.datasets import make_classification
from sklearn.model_selection import train_test_split, KFold, StratifiedKFold, GridSearchCV
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
import matplotlib.pyplot as plt
import seaborn as sns
```

2. Generate a Synthetic Dataset

Create a synthetic dataset with 2 classes

```
X, y = make_classification(
    n_samples=1000, n_features=10, n_informative=8, n_redundant=2,
    n_clusters_per_class=1, random_state=42
)
```

Convert to a DataFrame for visualization

```
df = pd.DataFrame(X, columns=[f'Feature_{i}' for i in range(1, 11)])
df['Target'] = y
```

Display the first few rows

```
print(df.head())
```

3. Split Data into Train and Test Sets

Split data into 80% training and 20% testing

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42,
                                                    stratify=y)
```

4. Define k-Fold Cross-Validation

```
kf = KFold(n_splits=5, shuffle=True, random_state=42)
print("k-Fold Cross-Validation:")
for train_index, val_index in kf.split(X_train):
    print("TRAIN:", train_index, "VALIDATION:", val_index)
```

5. Define Stratified k-Fold Cross-Validation

```
skf = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
print("\nStratified k-Fold Cross-Validation:")
for train_index, val_index in skf.split(X_train, y_train):
    print("TRAIN:", train_index, "VALIDATION:", val_index)
```

6. Train and Evaluate Using k-Fold Cross-Validation

Initialize model

```
model = RandomForestClassifier(random_state=42)
```

Perform k-Fold Cross-Validation

```

accuracies = []
for train_index, val_index in kf.split(X_train):
    X_kf_train, X_kf_val = X_train[train_index], X_train[val_index]
    y_kf_train, y_kf_val = y_train[train_index], y_train[val_index]

    # Train model
    model.fit(X_kf_train, y_kf_train)

    # Validate model
    y_pred = model.predict(X_kf_val)
    accuracy = accuracy_score(y_kf_val, y_pred)
    accuracies.append(accuracy)
print(f"Average Accuracy from k-Fold: {np.mean(accuracies):.2f}")

```

7. Hyperparameter Tuning Using GridSearchCV

Define parameter grid

```

param_grid = {
    'n_estimators': [50, 100, 200],
    'max_depth': [None, 10, 20, 30],
    'min_samples_split': [2, 5, 10],
}

```

Perform GridSearchCV with Stratified k-Fold

```

grid_search = GridSearchCV(
    estimator=RandomForestClassifier(random_state=42),
    param_grid=param_grid,
    cv=StratifiedKFold(n_splits=5, shuffle=True, random_state=42),
    scoring='accuracy',
    n_jobs=-1,
    verbose=1
)

```

Fit to training data

```

grid_search.fit(X_train, y_train)
print("Best Parameters:", grid_search.best_params_)
print("Best Cross-Validation Accuracy:", grid_search.best_score_)

```

8. Evaluate the Final Model

Use the best model for evaluation

```

best_model = grid_search.best_estimator_

```

Predict on test data

```

y_test_pred = best_model.predict(X_test)

```

Evaluate performance

```

print("\nTest Accuracy:", accuracy_score(y_test, y_test_pred))
print("\nClassification Report:\n", classification_report(y_test, y_test_pred))

```


Confusion matrix

```

conf_matrix = confusion_matrix(y_test, y_test_pred)
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', xticklabels=['Class 0', 'Class 1'],
            yticklabels=['Class 0', 'Class 1'])
plt.xlabel('Predicted')

```

```
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()
```



7b. Systematically explore combinations of hyperparameters to optimize model performance.(use grid and randomized search)

1. Import Necessary Libraries

```
import numpy as np
import pandas as pd
from sklearn.datasets import make_classification
from sklearn.model_selection import train_test_split, GridSearchCV, RandomizedSearchCV, StratifiedKFold
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
import matplotlib.pyplot as plt
import seaborn as sns
```

2. Generate a Synthetic Dataset

Generate a binary classification dataset

```
X, y = make_classification(
    n_samples=1000, n_features=12, n_informative=8, n_redundant=2,
    n_clusters_per_class=1, flip_y=0.03, random_state=42
)
```

Convert to a DataFrame for visualization

```
df = pd.DataFrame(X, columns=[f'Feature_{i}' for i in range(1, 13)])
df['Target'] = y
```

Display the first few rows

```
print(df.head())
```

3. Split Data into Train and Test Sets

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42, stratify=y)
```

4. Define the Model

```
# Initialize a Random Forest classifier
```

```
model = RandomForestClassifier(random_state=42)
```

5. Hyperparameter Tuning Using Grid Search

Define a parameter grid for Grid Search

```
param_grid = {
    'n_estimators': [50, 100, 200],
    'max_depth': [None, 10, 20],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4]
}
```

GridSearchCV with 5-fold cross-validation

```
grid_search = GridSearchCV(
    estimator=model,
    param_grid=param_grid,
    scoring='accuracy',
```

```

cv=StratifiedKFold(n_splits=5, shuffle=True, random_state=42),
verbose=1,
n_jobs=-1
)
# Fit the model
grid_search.fit(X_train, y_train)

# Best parameters and score from Grid Search
print("Best Parameters from Grid Search:", grid_search.best_params_)
print("Best Cross-Validation Accuracy from Grid Search:", grid_search.best_score_)

```

6. Hyperparameter Tuning Using Randomized Search

```

from scipy.stats import randint
# Define a parameter distribution for Randomized Search
param_dist = {
    'n_estimators': randint(50, 300),
    'max_depth': [None, 10, 20, 30],
    'min_samples_split': randint(2, 15),
    'min_samples_leaf': randint(1, 10)
}
# RandomizedSearchCV with 5-fold cross-validation
random_search = RandomizedSearchCV(
    estimator=model,
    param_distributions=param_dist,
    n_iter=50, # Number of random combinations to try
    scoring='accuracy',
    cv=StratifiedKFold(n_splits=5, shuffle=True, random_state=42),
    verbose=1,
    n_jobs=-1,
    random_state=42
)
# Fit the model
random_search.fit(X_train, y_train)

# Best parameters and score from Randomized Search
print("Best Parameters from Randomized Search:", random_search.best_params_)
print("Best Cross-Validation Accuracy from Randomized Search:",
random_search.best_score_)

```

7. Evaluate the Best Model

```

# Select the best model from Grid Search and Randomized Search
best_model = random_search.best_estimator_ # Or use grid_search.best_estimator_

# Predict on test data
y_test_pred = best_model.predict(X_test)

# Evaluate the performance
print("\nTest Accuracy:", accuracy_score(y_test, y_test_pred))
print("\nClassification Report:\n", classification_report(y_test, y_test_pred))

```


Confusion Matrix

```
conf_matrix = confusion_matrix(y_test, y_test_pred)
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', xticklabels=['Class 0', 'Class 1'],
            yticklabels=['Class 0', 'Class 1'])
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()
```

Test Accuracy: 0.94

Classification Report:

	precision	recall	f1-score	support
0	0.99	0.89	0.94	100
1	0.90	0.99	0.94	100
accuracy			0.94	200
macro avg	0.94	0.94	0.94	200
weighted avg	0.94	0.94	0.94	200



8. Implement Bayesian Learning using inferences

1. Import Necessary Libraries

```
import numpy as np
import pandas as pd
from sklearn.datasets import make_classification
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import GaussianNB
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt
```

2. Generate a Synthetic Dataset

Generate a dataset with 2 classes

```
X, y = make_classification(
    n_samples=1000, n_features=8, n_informative=6, n_redundant=2,
    n_classes=2, random_state=42
)
```

Convert to DataFrame for visualization

```
df = pd.DataFrame(X, columns=[f'Feature_{i}' for i in range(1, 9)])
df['Target'] = y
```

Display the first few rows

```
print(df.head())
```

3. Split the Dataset

Split data into 80% training and 20% testing

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42,
                                                    stratify=y)
```

4. Bayesian Learning with Naive Bayes

Initialize the Gaussian Naive Bayes model

```
model = GaussianNB()
```

Fit the model to the training data

```
model.fit(X_train, y_train)
```

Predict on the test data

```
y_pred = model.predict(X_test)
```

5. Evaluate the Model

Calculate accuracy

```
accuracy = accuracy_score(y_test, y_pred)
print(f"Test Accuracy: {accuracy:.2f}")
```

Print classification report

```
print("\nClassification Report:\n", classification_report(y_test, y_pred))
```

Generate and plot confusion matrix

```
conf_matrix = confusion_matrix(y_test, y_pred)
sns.heatmap(conf_matrix, annot=True, fmt='d', cmap='Blues', xticklabels=['Class 0', 'Class 1'],
yticklabels=['Class 0', 'Class 1'])
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix')
plt.show()
```

6. Understanding Bayesian Inference

In Bayesian Learning, the model predicts based on the probabilities:

- Prior Probability ($P(C)$): The likelihood of each class based on historical data.
- Likelihood ($P(X|C)$): The probability of the data given a class.
- Posterior Probability ($P(C|X)$): Calculated using Bayes' theorem:
$$P(C|X) = \frac{P(X|C) \cdot P(C)}{P(X)}$$

Example: Compute posterior probabilities for the first test sample

```
sample = X_test[0].reshape(1, -1)
posterior_probs = model.predict_proba(sample)
```

```
print(f"Sample Features: {sample}")
print(f"Posterior Probabilities: {posterior_probs}")
print(f"Predicted Class: {model.predict(sample)}")
```

